

Icosahedral Grid Graphs

We?

May 18, 2015

Revision History

Revision	Date	Author(s)	Description
-	01.10.2014	Carlos Osuna	Initial version

Contents

0.1	Introduction	5
0.2	Vertexes and Edges Graph	5
0.3	Domain Decomposition and Halos	10

0.1 Introduction

This document shows an schema for labelling the nodes of the graph of cells, vertexes and edges of an icosahedral grid. The grid is based on a R3B2 decomposition, where each edge of the original parallelogram of the icosahedral is tri-sect, and each of the resulting triangles is bi-sect (figure 1)




Figure 1: R3B2 decomposition of one of the parallelograms of the icosahedral. Blue lines show the first R3 decomposition, while black lines show the result of the additional B2 operation.

In this labelling, all the cells spin their upward, downward configuration in the grid, even when jumping to the next row. Least significant bit always identifies if the triangle is upward or downward. This layout of the blocks probably simplifies the looping along the cells when applying an operator. An example for the laplace follows

Figure also shows a possible way to tile the domain (in order to apply further parallelism) by always splitting into groups of two triangles with half size of the original parallelogram (figure 3)

```

1  // The following code shows only modifications to Mauro's triangular_storage
2  // that works on the graph labelling studied here for two triangles of one of
3  // the icosahedral parallelograms
4  struct triangular_offsets {
5      int m_offset[3];
6      static const int n_neighbors = 3;
7      triangular_offsets(int a) {
8          m_offset[0] = 1;
9          m_offset[1] = -1;
10         m_offset[2] = a;
11     }
12
13     int offset(int neighbor_index, int sign) const {
14         return sign*m_offset[neighbor_index];
15     }
16 };
17
18 template<typename OffsetFunction>
19 struct triangular_storage {
20
21     ...
22     iterator begin() {
23         return iterator(data.begin(), offset_function);
24     }
25     iterator operator++(int) const {
26         return m_it+1;
27     }
28     iterator operator+(int i) const {
29         return iterator(m_it+i, f, toggle_direction*(i&1 ? -1 : 1));
30     }
31     double& operator[](int i) {
32         return *(m_it+f.offset(i, toggle_direction);
33     }
34     template <typename Functor>
35     double fold_neighbors(iterator it, Functor && f) const {
36         double v = 0;
37         for (int i=0; i<OffsetFunction::n_neighbors; ++i) {
38             v = f(v, it[i]);
39         }
40         return v;
41     }
42 };
43
44 for (int i=1; i<n-1; ++i) {
45     for (int j=1; j<m-1; ++j) {
46         lap_cool.data[i*m+j] = 3*storage.data[i*m+j] -
47             storage.fold_neighbors(storage.begin()+i*m+j,
48                 [](double state, double value) {
49                     return state+value;
50                 });
51     }
52 }

```

Figure 2: Code example that illustrates the storage class and structure that defines the offset of neighbour cells, assuming an index layout like the one in figure 1

0.2 Vertexes and Edges Graph

In addition to the graph of center of cells, ICON will probably need to consider the graph of vertexes (figure 4) and edges (figure 5). Labeling of the nodes should allow to move along these graphs in a consistent way. A




Figure 3: One possible way of tiling an icosahedral parallelogram.

stencil operating on fields in the mass point (cell graph) might require access to the vertex graph storage or the edge graph storage. The data structure of the graph of cells could allow in the interface move into a different graph:

```

1 template<typename TEnv>
2 // structure that maps the cell graph into the vertex graph by
3 // providing the vertex ids around a cell id
4 //
5 // The VertexIdx is the index that identify one of the three vertexes that compose a triangle cell
6 // The three vertexes of a cell are numbered like (for downward triangle)
7 // (0)-----(1)
8 // \         /
9 //  \       /
10 //   \     /
11 //    \   /
12 //     \ /
13 //      (2)
14 // or for upward triangle
15 //
16 //      (0)
17 //     / \
18 //    /   \
19 //   /     \
20 //  /       \
21 // /         \
22 // (2)----- (1)
23 //
24 struct cell_map {
25     int m_offsets[6];
26     static const int n_vertexes = 3;
27     cell_map(int a) {
28         m_offset[0] = 0;
29         m_offset[1] = 1;
30         m_offset[2] = a/2;
31         m_offset[3] = 0;
32         m_offset[4] = a/2+1;
33         m_offset[5] = a/2;
34     }
35     // the following computes the vertex index of one of the vertexes
36     // (determined by VertexIdx) of a cell (CellIdx)

```

```

37 // Same idea should work with iterators
38 int vertex_index(const int VertexIdx, const int CellIdx){
39     return (CellIdx+1)/2+CellIdx/a + m_offset[VertexIdx+n_vertexes*(CellIdx&1)];
40 }
41 };
42
43 // structure that maps the cell graph into the edges graph by
44 // providing the edges ids around a cell id
45 //
46 // The EdgeIdx is the index that identify one of the three edges that compose a triangle cell
47 // The three edges of a cell are numbered like, for downward triangle
48 //
49 //   (1)
50 //  / \
51 // (0) (2)
52 //  \ /
53 //   \
54 //
55 //
56 // or for upward triangle
57 /*
58 //
59 //   / \
60 // (0) (1)
61 //  / \
62 // /---(2)---\
63 */
64
65 struct cell_map {
66     int m_offsets[6];
67     static const int n_edges = 3;
68     cell_map(int a) {
69         m_offset[0] = 0;
70         m_offset[1] = 1;
71         m_offset[2] = 2;
72         m_offset[3] = -1;
73         m_offset[4] = 0;
74         m_offset[5] = a/2*3-1;
75         // the following computes the edge index of one of the edges
76         // (determined by EdgeIdx) of a cell (CellIdx)
77     int edge_index(const int EdgeIdx, const int CellIdx){
78         return ((CellIdx+1)/2*3) + CellIdx/a + m_offset[EdgeIdx+n_edges*(CellIdx&1)];
79     }
80 };

```

Also we, similarly to `triangular_offsets`, each node of the vertex and edge graph should be able to (random) access and iterate over its neighbours. The following structure show how we can access the neighbours for the layout drawn in figure 4 and 5:

```

1 //
2 // indexes to refer to a vertex
3 //
4 //   0-----1
5 //  / \
6 // 5   ref   2
7 //  \ /
8 //   \
9 //   4-----3
10 //
11 //
12
13 struct vertex_neighbour_offsets {
14     int m_offsets[6];
15     static const int n_neighbours = 6;

```




Figure 4: Graph of icosahedral vertexes.

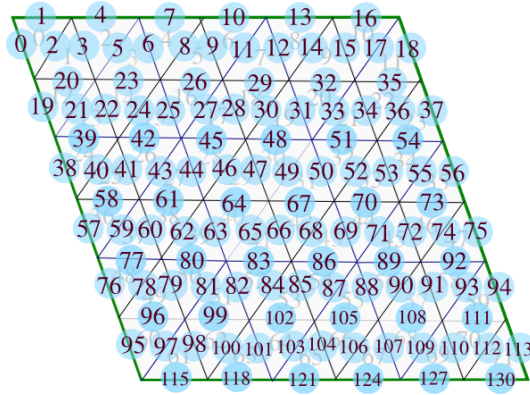


Figure 5: Graph of icosahedral edges. Note that the nodes in the last row show discontinuous index labelling in order to keep a consistent mapping between the graph of cells and the graph of edges. This requires adding some padding to the data structures.

```

16 vertex_neighbour_offsets(int a) {
17     m_offset[0] = -a/2-1;
18     m_offset[1] = -a/2;
19     m_offset[2] = 1;
20     m_offset[3] = a/2+1;
21     m_offset[4] = a/2;
22     m_offset[5] = -1;
23 }
24 // vertexes at the corners of the parallelogram will have only 5 neighbours
25 // i.e. vertexes 0, 6, 42, 48
26 // we will have to deal with these exceptional cases once we build halos around
27 // the parallelogram
28 int vertex_offset(int neighbour_index){
29     return m_offset[neighbour_index];
30 }
31 };

```

```

1 // There are three different shapes of parallelograms of edges with different
2 // neighbours relations. The neighbour index will is depicted below for each
3 // type of edge parallelogram
4 //
5 //     type I
6 //
7 //     1-----
8 //    / \
9 //   /   \
10 //  0   ref   2
11 // /     \
12 //-----3---
13 //
14 //
15 //     type II
16 //
17 //     1-----
18 //    / \
19 //   /   \
20 //  0   ref   2
21 // /     \
22 //-----3---
23 //
24 //
25 //     type III
26 //
27 //     / \
28 //    /   \
29 //   0     1
30 //  /       \
31 // --ref--
32 // /       \
33 //  3     2
34 //
35 //
36 //
37 //
38 struct edge_neighbour_offsets {
39     int m_offsets[4];
40     static const int n_neighbours = 4;
41     edge_neighbour_offsets(int a) {
42         m_offset[0] = -1;
43         m_offset[1] = 1;
44         m_offset[2] = 2;
45         m_offset[3] = a/2*3-1;
46         m_offset[4] = -2;
47         m_offset[5] = -1;
48         m_offset[6] = 1;
49         m_offset[7] = a/2*3;
50         m_offset[8] = -a/2*3;
51         m_offset[9] = -a/2*3;
52         m_offset[10] = 1;
53         m_offset[11] = -1;
54     }
55     int edge_offset(const int neighbour_index, const int type){
56         return m_offset[neighbour_index+type*n_neighbours];
57     }
58 };

```

0.3 Domain Decomposition and Halos

The original indexing proposed in figure 1 does not take into account halo regions, which will be essential in a block structured code. Adding halo regions around each parallelogram domain requires a well defined distribution of these parallelogram domains that fill the whole sphere.

Figure ?? shows an schema for numbering cells in a consistent way (similar to figure 1) but considering halo cells. Grid vertices at the corners of the diamond always have 5 surrounding cells, while all the other vertices in the grid contain 6 cells around.




Figure 6: Domain decomposition of a icosahedron grid into ten diamonds. Five of these contain the north pole, the other five the south pole

Indexing schema shown in figure ?? shows that all of the cells need the same offsets in order to access any of the three neighbouring cells (padding is required in order to achieve that). This is also true for most of the cells in the halos. However there are three cells around the corners of the diamond, enclosed in blue circles in figure ??, which require different offsets. I.e. the triangular offsets defined in figure 2 is not valid for those cells

under blue circles. These grid cells will require special treatment.

This can also impose some restrictions, also on cells inside the computation domain in case operators of order greater than one are used¹ For example if second order operators are defined by the user, i.e. stencils that access second order neighbours, the corner cells 17, 28 and 98 will also require special treatment.

In case one of the original diamonds in figure 9 is further decomposed, the corner vertices of the resulting diamond is always composed by six neighbour cells (see figure 10). In this case, the offset pattern of neighbour is always applicable for all the grid cells (included halo cells).

¹This refer to high order operators which are computed on the fly. For example a fourth order horizontal diffusion can be computed as two laplace operators concatenated via a temporary field, or as two nested laplace stencil functions. The first case should never be a limitation here, as each operator will only access first order neighbours.




Figure 7: Grid cells labeling of one of the original diamonds containing halo regions with 2 neighbour cells at the halos. Red triangle represents padding cells required to keep consistency of strides along the different rows of the diamond. Each diamond of the original icosahedron decomposition is identified with the label convention followed in figure 9. Arrows indicate neighbour relation between two cells which are not contiguous in the indexing space (there is some padding). The blue wedge indicates a transition between two regions that do not respect the neighborhood rules of cells. Yellow wedge indicates a transition between triangles where contiguous downward/upward swap is not respected (i.e. triangles 1 and 2 are both upward). Special care must be taken when jumping into a padding cell, following any of the arrows connecting two cells and/or crossing any of the two wedges. In that case the neighborhood rules dictated by `triangular_offsets` in Figure 2 are not respected.

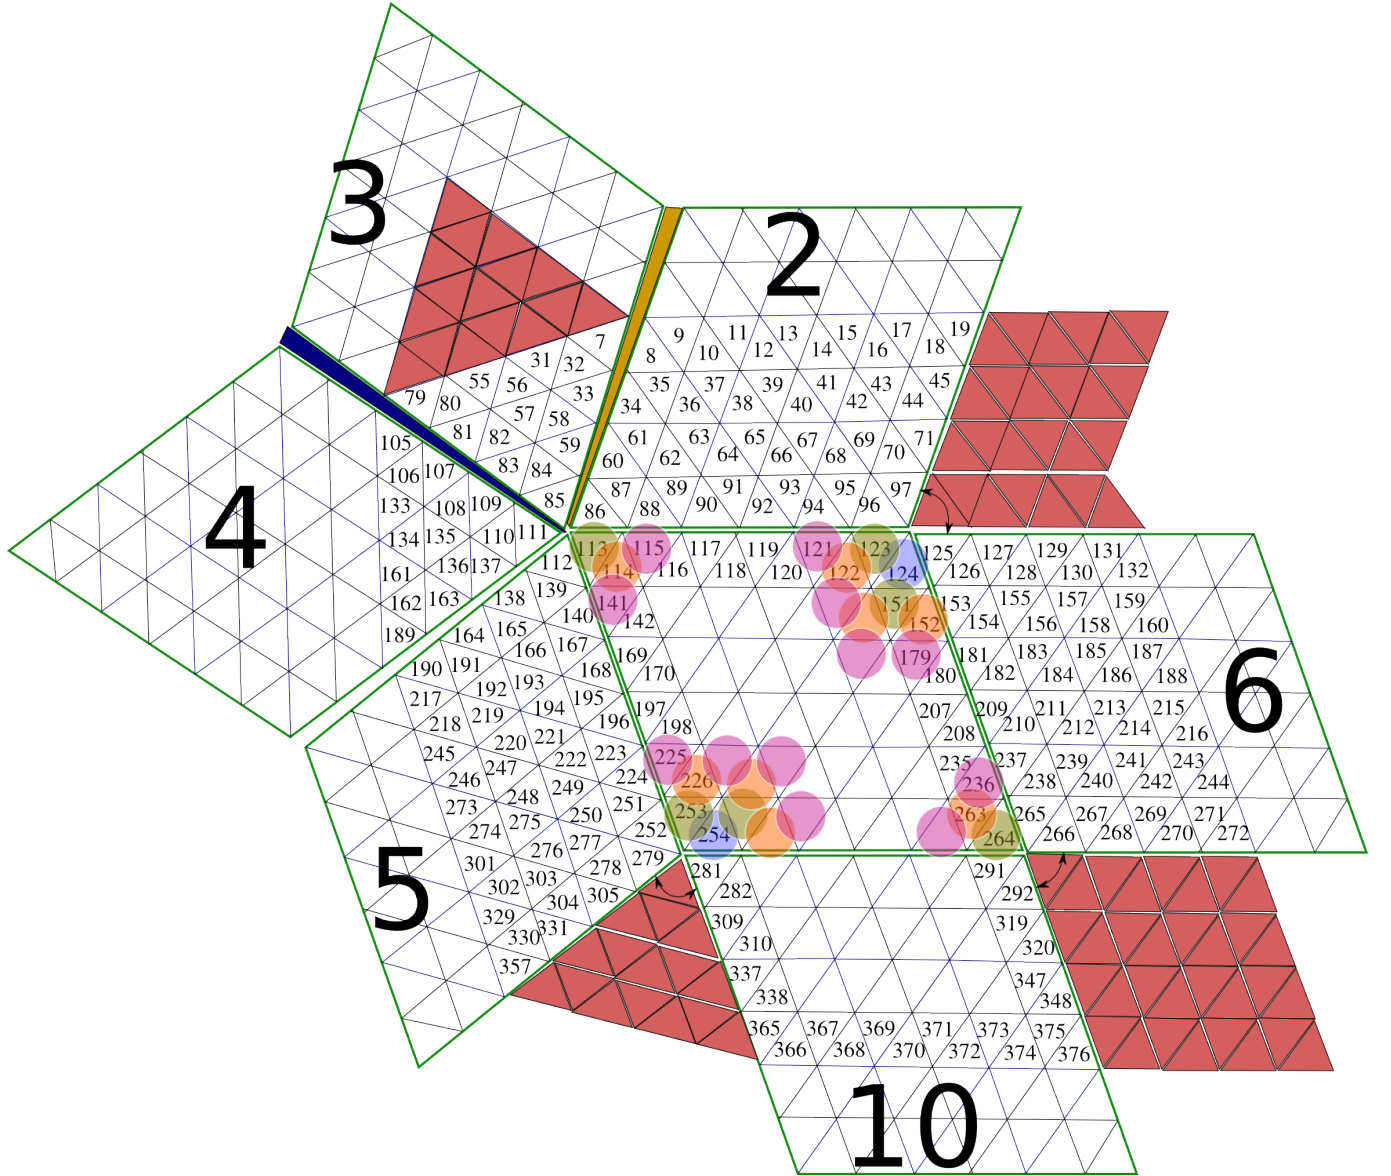


Figure 8: Same as Figure 7 with 4 halo lines. Red triangle represents padding cells required to keep consistency of strides along the different rows of the diamond. Each diamond of the original icosahedron decomposition is identified with the label convention followed in figure 9. Arrows indicate neighbour relation between two cells which are not contiguous in the indexing space (there is some padding). The blue wedge indicates a transition between two regions that do not respect the neighborhood rules of cells. Yellow wedge indicates a transition between triangles where contiguous downward/upward swap is not respected (i.e. triangles 1 and 2 are both upward). Special care must be taken when jumping into a padding cell, following any of the arrows connecting two cells and/or crossing any of the two wedges. In that case the neighborhood rules dictated by `triangular_offsets` in Figure 2 are not respected. None of the cells in the computation domain can cross any of these regions (or fall into a padding cell) with one neighbour move. Blue circles mark cells which can cross one of these regions (or fall into a padding cell) in which case the general offset rule is not respected with two neighbour moves. Green circles need three neighbour moves. Orange circles need 4 neighbour moves and finally magenta need 5.




Figure 9: Domain decomposition of an icosahedron grid into ten diamonds. Five of these contain the north pole, the other five the south pole




Figure 10: Grids cells labeling of an inner diamond, where none of the corners is a corner of a diamond of the original decomposition of the icosahedron (figure 9). Red cells indicate required padding in order to keep uniform offset rules accross the domain. In this case, all the cells exhibit exactly the same neighbour cells relationship (i.e. describe with a unique offset for the whole domain and halos). No exceptional treatment of some cells is required.